

Transfer learning using pretrained ImageNet models has been the standard approach for chest X ray tasks and many other medical imaging modalities. ImageNet pretraining yields a statistically significant boost in performance across architectures, with a higher boost for smaller architectures. The truncated models maintained favorable accuracy to parameter tradeoffs, making compact backbones attractive for resource constrained medical imaging deployments.